Task :

Create a flume agent that streams data from Twitter and stores in the HDFS.

Go to Twitter account and create and application on Twitter apps

Save below Consumer Key, Consumer secret , Access Token, Access Token Secret and update into configuration file.

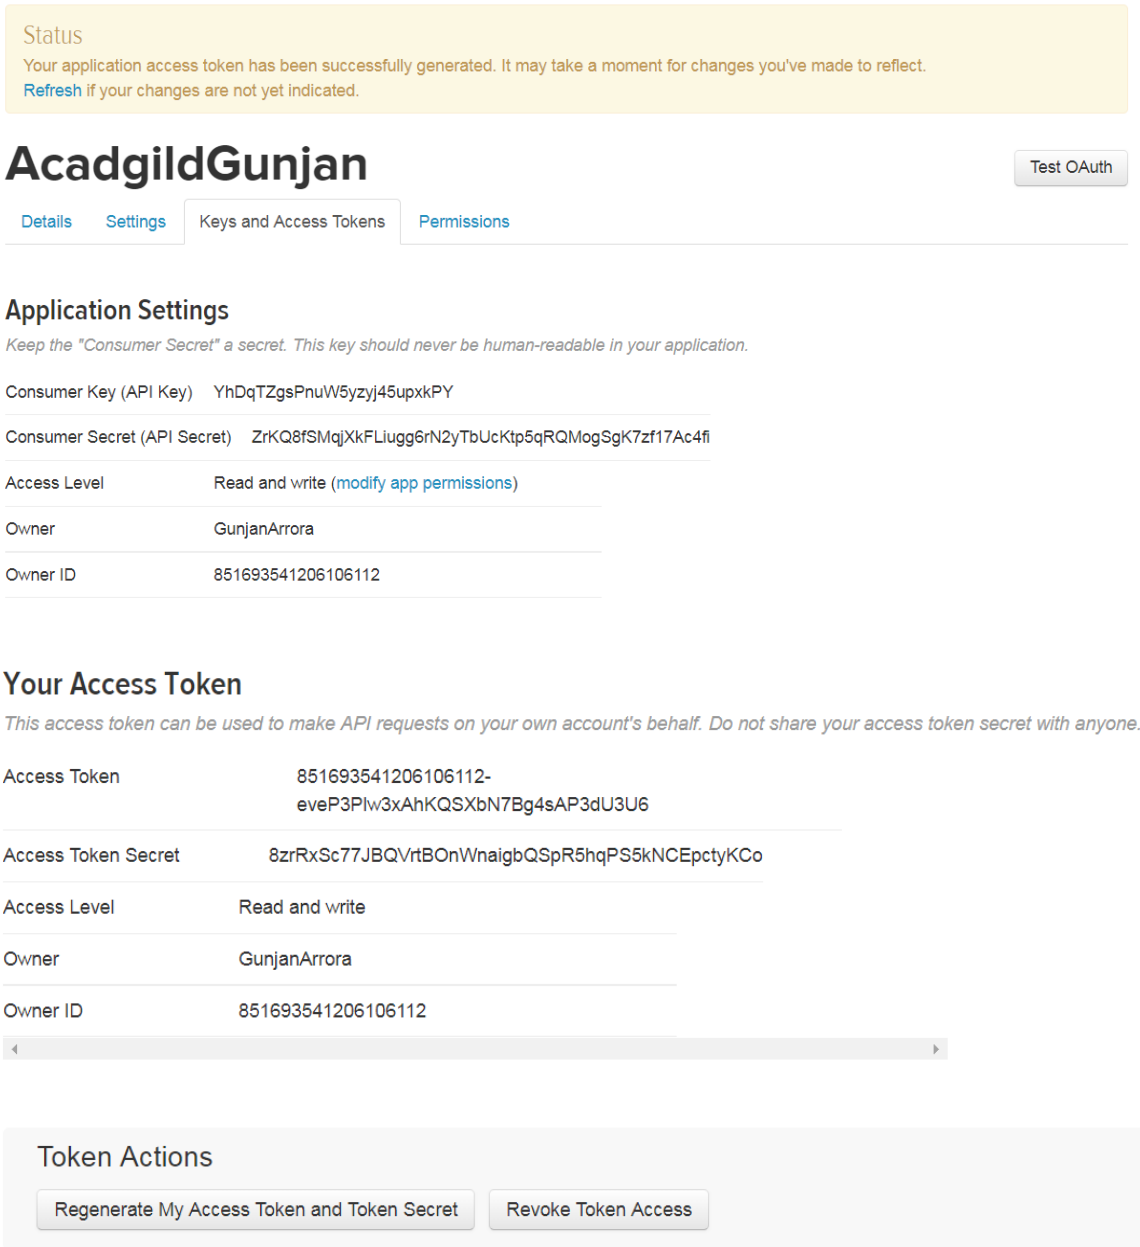

➔ Check all jar files of twitter, weather they are available or not

Tp make sure below jars placed in $FLUME_HOME/lib directory:

- ➔ *twitter4j-core-X.XX.jar*
- ➔ *twitter4j-stream-X.X.X.jar*
- ➔ *twitter4j-media-support-X.X.X.jar*

Configuration file: (Update Configuration file with latest key – values and put into VM-Flume path

```
TwitterAgent.sources = Twitter
TwitterAgent.channels = MemChannel
TwitterAgent.sinks = HDFS

# Describing/Configuring the source
TwitterAgent.sources.Twitter.type =
org.apache.flume.source.twitter.TwitterSource
TwitterAgent.sources.Twitter.consumerKey=YhDqTZgsPnuW5yzyj45upxkP
Y
TwitterAgent.sources.Twitter.consumerSecret=ZrKQ8fSMqjXkFLiugg6rN
2yTbUcKtp5qRQMogSgK7zf17Ac4fi
TwitterAgent.sources.Twitter.accessToken=851693541206106112-
eveP3Plw3xAhKQSXbN7Bg4sAP3dU3U6
TwitterAgent.sources.Twitter.accessTokenSecret=
8zrRxSc77JBQVrtBOnWnaigbQSpR5hqPS5kNCEpctyKCo
TwitterAgent.sources.Twitter.keywords=hadoop, bigdata, mapreduce,
mahout, hbase, nosql
# Describing/Configuring the sink

TwitterAgent.sources.Twitter.keywords= hadoop,election,sports,
cricket,Big data,gunjan,example

TwitterAgent.sinks.HDFS.channel=MemChannel
TwitterAgent.sinks.HDFS.type=hdfs
TwitterAgent.sinks.HDFS.hdfs.path=hdfs://localhost:9000/user/flum
e/tweets
TwitterAgent.sinks.HDFS.hdfs.fileType=DataStream
TwitterAgent.sinks.HDFS.hdfs.writeformat=Text
TwitterAgent.sinks.HDFS.hdfs.batchSize=1000
TwitterAgent.sinks.HDFS.hdfs.rollSize=0
TwitterAgent.sinks.HDFS.hdfs.rollCount=10000
TwitterAgent.sinks.HDFS.hdfs.rollInterval=600

TwitterAgent.channels.MemChannel.type=memory
TwitterAgent.channels.MemChannel.capacity=10000
TwitterAgent.channels.MemChannel.transactionCapacity=1000

TwitterAgent.sources.Twitter.channels = MemChannel
TwitterAgent.sinks.HDFS.channel = MemChannel
```

Create directory on hdfs with names tweets.

```
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -mkdir -p /user/flume/tweets/
```

```
[acadgild@localhost ~]$ hadoop fs -ls /user/flume
18/05/05 10:22:03 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platform... using builtin-java classes where applicable
Found 1 items
drwxr-xr-x   - acadgild supergroup          0 2018-05-05 10:20 /user/flume/tweet
s
```

For fetching data from Twitter, Use the below command to fetch the twitter tweet data into the HDFS cluster path.

```
[acadgild@localhost ~]$ flume-ng agent -n TwitterAgent -f /home/acadgild/install
/flume/apache-flume-1.8.0-bin/conf/acadgildTwitterFlume.conf
```

The above command will start fetching data from Twitter and steams it into the HDFS given path.

Streaming:

```
, code=-1, retryAfter=-1, rateLimitStatus=null, version=3.0.3}
        at twitter4j.internal.http.HttpClientImpl.request(HttpClientImpl.java:17
7)
        at twitter4j.internal.http.HttpClientWrapper.request(HttpClientWrapper.j
ava:61)
        at twitter4j.internal.http.HttpClientWrapper.get(HttpClientWrapper.java:
89)
        at twitter4j.TwitterStreamImpl.getSampleStream(TwitterStreamImpl.java:17
6)
        at twitter4j.TwitterStreamImpl$4.getStream(TwitterStreamImpl.java:164)
        at twitter4j.TwitterStreamImpl$TwitterStreamConsumer.run(TwitterStreamIm
pl.java:462)
18/05/05 10:26:57 INFO twitter4j.TwitterStreamImpl: Waiting for 160000 milliseco
nds
```

```
twitter.TwitterSource: Processed 100 docs
twitter.TwitterSource: Processed 200 docs
twitter.TwitterSource: Processed 300 docs
twitter.TwitterSource: Processed 400 docs
twitter.TwitterSource: Processed 500 docs
```

Once, the tweet data started streaming it into the given HDFS path we can use 'Ctrl+c' command to stop the streaming process

- ⇨ To check the contents of the tweet data we can use the following command:**hadoop fs –ls /user/flume/tweets**

```
INFO hdfs.BucketWriter: Creating hdfs://localhost:9000/user/flume/tweets/FlumeData.1523160169655.
```

Use the 'cat' command to display the tweet data inside the /user/flume/tweets/FlumeData.1523160169655 path using

hadoop dfs –cat /user/flume/tweets **/FlumeData.1523160169655**